Mixed Content Evaluation Document

Abstract

This document combines tables, equations, figures, and bibliographic references in a single page to stress-test multi-element extraction. Section 1 presents a results table. Section 2 introduces Equation 1. The conclusion references Figure 1.

Table 1. Extraction Results by Strategy

Extractor	Recall	Precision	F1	Latency
kreuzberg	0.91	0.88	0.895	120ms
docling	0.86	0.84	0.850	380ms
pymupdf4llm	0.79	0.82	0.805	90ms
pdftotext	0.70	0.75	0.724	40ms

Section 2: Evaluation Metric

```
Equation 1: F1 = 2 * (precision * recall) / (precision + recall)
```

The F1 score is the harmonic mean of precision and recall, providing a balanced assessment when the class distribution is uneven. See reference [1] for derivation.

Conclusion

Figure 1 and Table 1 together demonstrate that kreuzberg achieves the best F1 score while maintaining competitive latency. Future work will extend Eq. 1 to a weighted variant. See reference [2] for the full dataset.

References

[1] Manning et al. (2008). Introduction to Information Retrieval.

[2] DocSlurp Team (2025). Kreuzberg Eval Dataset. https://data.docslurp.io/evals